


```
In [164]: etprediction = model2.predict(X_test)

In [165]: etprediction

Out [65]: array([0.17, 0., 0., 0., 0., 0., 0.03, 3.21, 0.8, 0., 1.73, 2.62,
1.28, 0.02, 0., 0., 0.01, 0.39, 0.36, 0., 0.46, 0., 0., 0.03,
1.16, 0.28, 0.6, 0., 0., 2.32, 0.46, 0.05, 1.96, 0., 0., 0.03,
0.11, 0.02, 0., 0., 0., 2.07, 0., 0.11, 0.03, 0.23, 0.08, 0.,
0., 0., 0., 0., 0.29, 0., 0.13, 0.35, 0.72, 0.08, 0.38,
0., 2.5, 0., 0.04, 0., 0., 0.04, 0., 0., 0.16, 1.46,
0., 0.05, 0.44, 0., 1.54, 0.65, 0., 0., 1.79, 2., 1.89,
0.71])

In [166]: etprediction.round()

Out [66]: array([0., 0., 0., 0., 0., 0., 0., 3., 1., 0., 2., 3., 1., 0., 0., 0., 0.,
0., 0., 0., 0., 1., 0., 0., 0., 0., 2., 0., 0., 0., 0., 0., 0.,
4., 1., 0., 0., 0., 2., 0., 0., 0., 0., 0., 0., 1., 1., 0., 0.,
0., 0., 2., 1., 0., 0., 0., 0., 2., 2., 2., 1.]

Model Evaluation

In [167]: mse = mean_squared_error(y_test,etprediction.round())
mse

Out [67]: 2.282051282051282

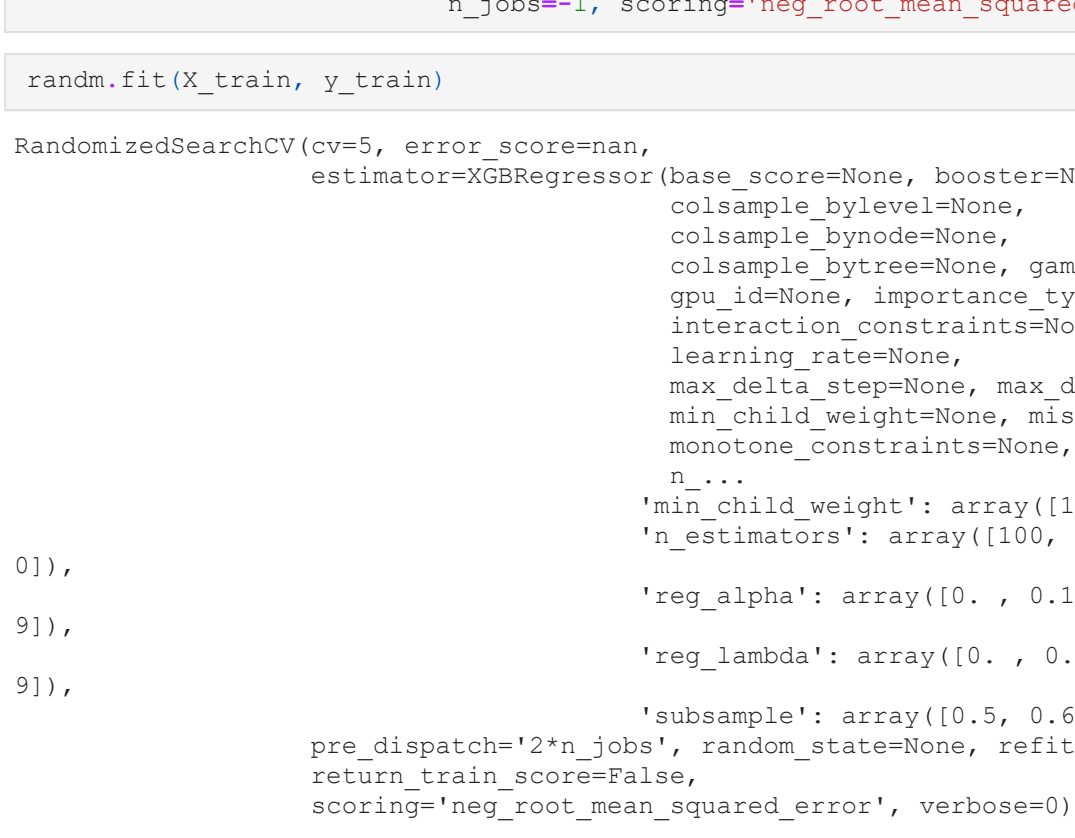
In [168]: rmse = np.sqrt(mse)

Out [68]: 1.5106459817082498

In [169]: r2score = r2_score(y_test,etprediction.round())

Out [69]: 0.2378129117259552

In [170]: fig, ax = plt.subplots(figsize=(10,4))
sns.regplot(x=y_test, y=etprediction.round(), ax=ax)
plt.title("Plot to compare actual vs predicted - Extra Trees Regressor")
plt.ylabel("Predicted")
plt.xlabel("Actual")
plt.show()
```



Using XGBoost (Scikit-Learn)

Using RandomSearchCV

```
In [171]: model3 = XGBRegressor(random_state=0, n_estimators=100, objective='reg:squarederror')

In [172]: parameters = {'max_depth': np.arange(3,10,1),
                        'eta': np.arange(0.05,0.3,0.05),
                        'n_estimators': np.arange(100,1000,100),
                        'min_child_weight': np.arange(1,4,1),
                        'gamma': np.arange(0,10,2),
                        'subsample': np.arange(0.5,0.9,0.1),
                        'colsample_bytree': np.arange(0.5,0.9,0.1),
                        'reg_alpha': np.arange(0,1,0.1),
                        'reg_lambda': np.arange(0,1,0.1)
                        }

In [173]: randm = RandomizedSearchCV(estimator=model3, param_distributions = parameters, cv = 5, n_iter = 20,
                                  n_jobs=-1, scoring='neg_root_mean_squared_error')

In [174]: randm.fit(X_train, y_train)

Out [174]: RandomizedSearchCV(cv=5, error_score=nan,
                             estimator=XGBRegressor(base_score=None, booster=None,
                             colsample_bytree=None,
                             colsample_bynode=None,
                             colsample_bylevel=None,
                             gamma=None,
                             gpu_id=None, importance_type='gain',
                             interaction_constraints=None,
                             learning_rate=None,
                             max_delta_step=None,
                             max_depth=None,
                             min_child_weight=None,
                             missing=None,
                             monotone_constraints=None,
                             n_estimators=100,
                             n_jobs=-1,
                             num_parallel_tree=None,
                             random_state=None,
                             refit=True,
                             reg_alpha=0.0,
                             reg_lambda=0.0,
                             scale_pos_weight=1.0,
                             subsample=1.0,
                             tree_method='exact',
                             validate_parameters=1,
                             verbosity=None)

In [175]: randm.best_estimator_

Out [175]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bylevel=1, eta=0.05, gamma=4,
                        gpu_id=-1, importance_type='gain', interaction_constraints='',
                        learning_rate=0.050000007, max_delta_step=0, max_depth=5,
                        min_child_weight=1, missing=None, monotone_constraints=(),
                        n_estimators=400, n_jobs=0, num_parallel_tree=1,
                        objective='reg:squarederror', random_state=0,
                        reg_alpha=0.7000000000000001, reg_lambda=0.30000000000000004,
                        scale_pos_weight=1, subsample=0.5, tree_method='exact',
                        validate_parameters=1, verbosity=None)

In [176]: randm.best_score_

Out [176]: -1.3828418002128893

In [177]: randm.best_params_

Out [177]: {'subsample': 0.5,
'eta': 0.30000000000000004,
'reg_alpha': 0.7000000000000001,
'n_estimators': 400,
'min_child_weight': 3,
'max_depth': 5,
'gamma': 4,
'colsample_bytree': 0.6}
```

Final Model

```
In [178]: xgbmodel = XGBRegressor(random_state=0, n_estimators=400, objective='reg:squarederror',
                                subsample=0.5,reg_lambda=0.3,reg_alpha=0.7,min_child_weight=3,max_depth=5,
                                gamma=4, eta=0.05,colsample_bytree=0.6)

In [179]: xgbmodel.fit(X_train,y_train,eval_set=(X_test,y_test),eval_metric='rmse',early_stopping_rounds=10)

Will train until validation_0-rmse hasn't improved in 10 rounds.
[1] validation_0-rmse:1.75821
[2] validation_0-rmse:1.68752
[3] validation_0-rmse:1.64542
[4] validation_0-rmse:1.62053
[5] validation_0-rmse:1.58318
[6] validation_0-rmse:1.56379
[7] validation_0-rmse:1.54993
[8] validation_0-rmse:1.52237
[9] validation_0-rmse:1.49889
[10] validation_0-rmse:1.47791
[11] validation_0-rmse:1.46083
[12] validation_0-rmse:1.43904
[13] validation_0-rmse:1.43089
[14] validation_0-rmse:1.41257
[15] validation_0-rmse:1.39626
[16] validation_0-rmse:1.38596
[17] validation_0-rmse:1.38386
[18] validation_0-rmse:1.37286
[19] validation_0-rmse:1.36382
[20] validation_0-rmse:1.35936
[21] validation_0-rmse:1.35815
[22] validation_0-rmse:1.35558
[23] validation_0-rmse:1.35107
[24] validation_0-rmse:1.35303
[25] validation_0-rmse:1.34999
[26] validation_0-rmse:1.33919
[27] validation_0-rmse:1.33155
[28] validation_0-rmse:1.33065
[29] validation_0-rmse:1.33051
[30] validation_0-rmse:1.33587
[31] validation_0-rmse:1.33156
[32] validation_0-rmse:1.32577
[33] validation_0-rmse:1.32992
[34] validation_0-rmse:1.33098
[35] validation_0-rmse:1.32998
[36] validation_0-rmse:1.33051
[37] validation_0-rmse:1.33100
[38] validation_0-rmse:1.32762
[39] validation_0-rmse:1.31356
[40] validation_0-rmse:1.31411
[41] validation_0-rmse:1.31607
[42] validation_0-rmse:1.30938
[43] validation_0-rmse:1.30557
[44] validation_0-rmse:1.30499
[45] validation_0-rmse:1.30563
[46] validation_0-rmse:1.30046
[47] validation_0-rmse:1.29726
[48] validation_0-rmse:1.29286
[49] validation_0-rmse:1.29407
[50] validation_0-rmse:1.29202
[51] validation_0-rmse:1.29470
[52] validation_0-rmse:1.29148
[53] validation_0-rmse:1.28896
[54] validation_0-rmse:1.28999
[55] validation_0-rmse:1.28949
[56] validation_0-rmse:1.28967
[57] validation_0-rmse:1.28944
[58] validation_0-rmse:1.28951
[59] validation_0-rmse:1.28778
[60] validation_0-rmse:1.28749
[61] validation_0-rmse:1.29161
[62] validation_0-rmse:1.28673
[63] validation_0-rmse:1.28983
[64] validation_0-rmse:1.29113
[65] validation_0-rmse:1.29351
[66] validation_0-rmse:1.29351
[67] validation_0-rmse:1.30079
[68] validation_0-rmse:1.29351
[69] validation_0-rmse:1.30075
[70] validation_0-rmse:1.29939
[71] validation_0-rmse:1.29410
[72] validation_0-rmse:1.29620
Stopping. Best iteration!
[65] validation_0-rmse:1.28693

Out [179]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bylevel=1, eta=0.05, gamma=4,
                        gpu_id=-1, importance_type='gain', interaction_constraints='',
                        learning_rate=0.050000007, max_delta_step=0, max_depth=5,
                        min_child_weight=1, missing=None, monotone_constraints=(),
                        n_estimators=400, n_jobs=0, num_parallel_tree=1,
                        objective='reg:squarederror', random_state=0,
                        reg_alpha=0.7,
                        reg_lambda=0.3, scale_pos_weight=1, subsample=0.5,
                        tree_method='exact', validate_parameters=1, verbosity=None)

In [180]: y_pred = xgbmodel.predict(X_test)

In [181]: y_pred

Out [181]: array([1.693808, 0.13319057, 0.05102912, 0.20768976, 0.20737144,
0.24536291, 2.4978304, 0.386319, 0.00288917, 0.09225506, 0.06289913,
0.00852917, 0.17013602, 0.01083846, 0.43664845, 0.07461962,
0.43183863, 0.20026293, 1.6112388, 0.02448371, 0.00714466, 0.00274756,
0.00176685, 0.00506808, 0.01110173, 0.00147041, 0.00071154,
0.00750016, 0.0058574, 0.00214861, 0.00833984, 0.00454469,
0.13319057, 0.14977062, 2.1358585, 0.3020173, 0.31540662,
0.12845254, 0.08497405, 4.5722485, 0.15236435, 0.22618972,
0.30875778, 2.0225163, 0.1952821, 0.07461962, 0.13319057,
0.13319057, 0.11525298, 0.19766959, 1.8202349, 0.12877086,
0.20026293, 3.6301598, 0.46931675, 0.11331057, 1.8289408,
0.08497405, 3.348646, 0.07461962, 0.586391, 0.07461962,
0.0746362, 0.5962136, 0.25558794, 0.18085134, 0.8977414,
1.0813393, 0.09854544, 0.15654445, 1.6898073, 0.08487405,
2.3818988, 0.1284254, 0.35158795, 0.03839973, 0.00377258,
0.09324342, 0.00740652, 0.02676879, 0.0381528, 0.02985687,
0.2832831, 0.15066561])

In [182]: y_pred.round()

Out [182]: array([2., 0., 0., 0., 0., 2., 1., 0., 3., 3., 1., 0., 0., 0., 0.,
2., 0., 2., 0., 0., 1., 1., 0., 0., 2., 0., 0., 0., 0., 2., 0.,
0., 0., 0., 2., 0., 3., 0., 1., 0., 0., 0., 0., 2., 0., 0.,
4., 1., 0., 0., 0., 3., 0., 1., 0., 0., 0., 1., 0., 1., 1., 0.,
2., 0., 2., 0., 0., 0., 3., 4., 3., 1.], dtype=float32)
```

Plot Feature Importances (Extra Tree and XGBoost)

```
In [187]: et.feature_importances_

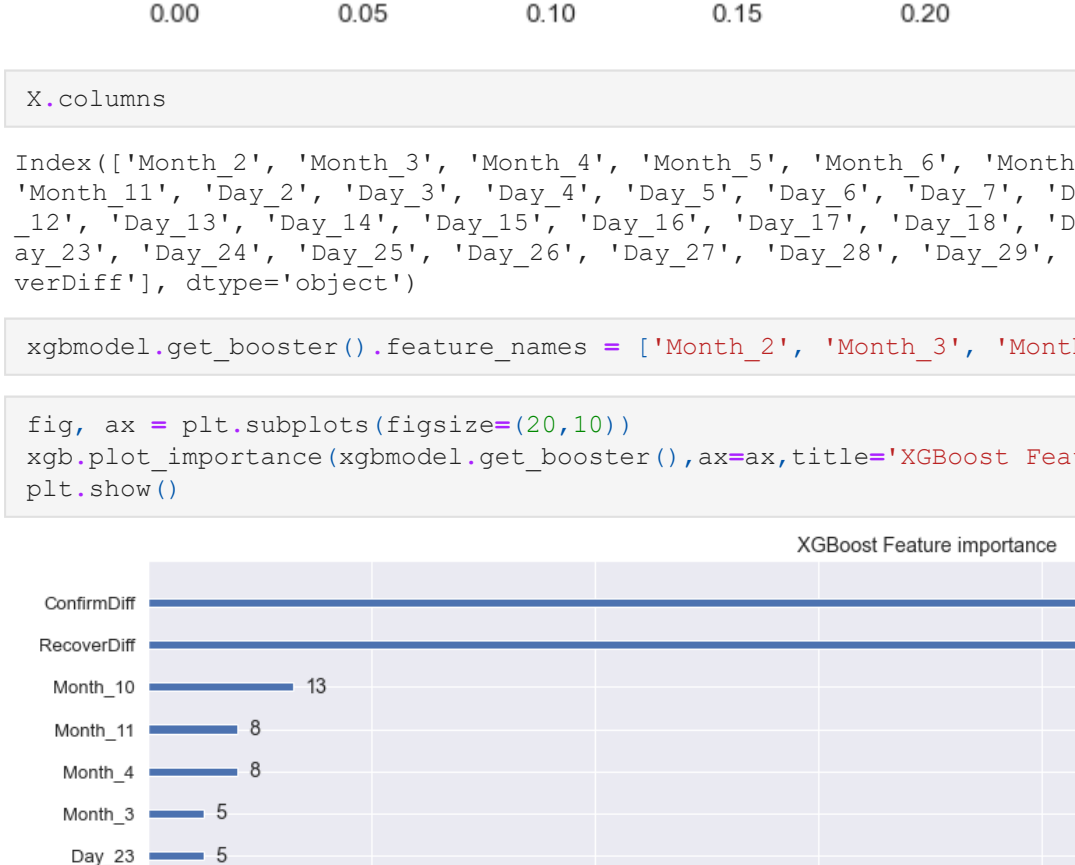
Out [187]: array([0.00002371, 0.01385562, 0.03296161, 0.00464944, 0.00149913,
0.00098501, 0.00020326, 0.00288917, 0.09225506, 0.06289913,
0.00852917, 0.17013602, 0.01083846, 0.43664845, 0.07461962,
0.00315433, 0.01444801, 0.02448371, 0.00714466, 0.00274756,
0.00176685, 0.00506808, 0.01110173, 0.00147041, 0.00071154,
0.00750016, 0.0058574, 0.00214861, 0.00833984, 0.00454469,
0.00369115, 0.04312089, 0.00354092, 0.00839973, 0.00377258,
0.00324342, 0.00740652, 0.02676879, 0.0381528, 0.02985687,
0.2832831, 0.15066561])

In [188]: feat_importances = pd.Series(et.feature_importances_, index=X.columns)

In [189]: feat_importances

Out [189]: Month_2    0.000024
Month_3    0.013856
Month_4    0.032962
Month_5    0.004649
Month_6    0.001499
Month_7    0.000985
Month_8    0.002033
Month_9    0.002889
Month_10   0.092255
Month_11   0.062899
Day_2     0.008282
Day_3     0.001377
Day_4     0.023202
Day_5     0.012272
Day_6     0.000754
Day_7     0.003154
Day_8     0.014448
Day_9     0.044884
Day_10    0.007145
Day_11    0.002748
Day_12    0.001767
Day_13    0.003068
Day_14    0.011102
Day_15    0.001470
Day_16    0.000712
Day_17    0.007501
Day_18    0.005858
Day_19    0.002149
Day_20    0.008340
Day_21    0.004545
Day_22    0.003691
Day_23    0.043121
Day_24    0.003541
Day_25    0.008399
Day_26    0.003773
Day_27    0.002434
Day_28    0.007407
Day_29    0.026769
Day_30    0.038153
Day_31    0.029857
ConfirmDiff 0.282828
RecoverDiff 0.150267
dtype: float64

In [190]: feat_importances.nlargest(10).plot(kind='barh', figsize=(10,10))
plt.title("Extra Tree Feature Importances")
plt.show()
```



```
In [191]: X.columns

Out [191]: Index(['Month_2', 'Month_3', 'Month_4', 'Month_5', 'Month_6', 'Month_7', 'Month_8', 'Month_9', 'Month_10',
'Month_11', 'Day_2', 'Day_3', 'Day_4', 'Day_5', 'Day_6', 'Day_7', 'Day_8', 'Day_9', 'Day_10', 'Day_11', 'Day_12',
'Day_13', 'Day_14', 'Day_15', 'Day_16', 'Day_17', 'Day_18', 'Day_19', 'Day_20', 'Day_21', 'Day_22', 'Day_23',
'Day_24', 'Day_25', 'Day_26', 'Day_27', 'Day_28', 'Day_29', 'Day_30', 'Day_31', 'ConfirmDiff', 'RecoverDiff'],
dtype='object')

In [192]: xgbmodel.get_booster().feature_names = ['Month_2', 'Month_3', 'Month_4', 'Month_5', 'Month_6', 'Month_7', 'M
```

```
In [193]: fig, ax = plt.subplots(figsize=(20,10))
xgb.plot_importance(xgbmodel.get_booster(), ax=ax, title='XGBoost Feature importance')
plt.show()
```



Cross-Validation

```
In [194]: cv = cross_val_score(xgbmodel, X, y, cv=5, verbose=1, scoring='neg_root_mean_squared_error')

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 5 out of 5 | elapsed: 1.1s finished

Out [195]: -1.4765106992218895

XGBoost Model Prediction on Test Dataset

In [198]: testref = pd.read_csv('testccovid.csv')

In [199]: testref

Out [199]:
```

	Month_2	Month_3	Month_4	Month_5	Month_6	Month_7	Month_8	Month_9	Month_10	Month_11	...	Day_25	Day_26	Day_27	Da
0	0	0	0	0	0	0	0	0	0	0	1	...	0	1	0
1	0	0	0	0	0	0	0	0	0	0	1	...	0	0	1
2	0	0	0	0	0	0	0	0	0	0	1	...	0	0	0

3 rows x 43 columns

```
In [100]: testdata = pd.read_csv('testccovid2.csv')

In [101]: testdata

Out [101]:
```

	Month_2	Month_3	Month_4	Month_5	Month_6	Month_7	Month_8	Month_9	Month_10	Month_11	...	Day_24	Day_25	Day_26	Da
0	0	0	0	0	0	0	0	0	0	0	1	...	0	0	1
1	0	0	0	0	0	0	0	0	0	0	1	...	0	0	0
2	0	0	0	0	0	0	0	0	0	0	1	...	0	0	0

3 rows x 42 columns

```
In [102]: testprediction = xgbmodel.predict(testdata)

In [104]: testprediction

Out [104]: array([2.6429937, 3.0189664, 2.5625987], dtype=float32)
```

```
In [105]: testprediction.round()

Out [105]: array([3., 3., 3.], dtype=float32)
```

```
In [106]: testref['DeathsDiff']

Out [106]: 0    3.0
1    1.0
2    4.0
Name: DeathsDiff, dtype: float64

Save the Model

In [107]: filename = 'xgbmodel(MN).sav'
dump(xgbmodel, open(filename, 'wb'))

In [ ]:
```